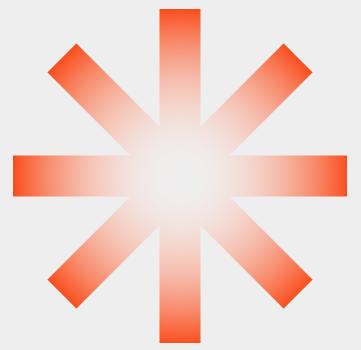


Directed k-MTM Multicast Implementation

Approximation Algorithms in the Telephone Model

PRESENTED BY

Hamza Abdullah



Overview

Problem & Motivation

Results

Algorithm

Demo

Implementation

Conclusion

Problem & Motivation

- Objective: Inform any $k \leq t$ terminals in a directed network under the telephone model
 - Each informed node may call one neighbor per synchronous round
- Why it matters:
 - Database replication consistency
 - Sensor-network firmware updates
 - Vector-clock synchronization delay
- Core challenge:
 - NP-hard combinatorial problem
 - Matching constraint (one call/node/round) limits throughput

Algorithm

- Greedy Packing
 - Extract $\rho = \text{sq_root}[k]$ vertex-disjoint subtrees of depth $\leq D^*$, each covering ρ terminals.
- Partition Matroid Cover
 - Cover the remaining terminals
 - Half-approx greedy
 - Continuous-greedy
 - Lazy-greedy speed-up
- Stitch & Simulate
 - Combine subtrees + cover edges into one tree, then simulate telephone-model rounds to count broadcast time.

Implementation

- `graph_loader.py` – generates ER and clique graphs, picks ttt and kkk ratios
- `greedy_packing.py` – extracts the depth- D^* subtrees
- `pmcover.py`, `pmcover_continuous.py`, `pmcover_lazy.py` – three variants of matroid-constrained set-cover
- `complete.py` – stitches subtrees and cover edges into a single multicast tree
- `simulator.py` – runs the telephone-model broadcast rounds on the final tree

Enhancements:

- Lazy-greedy speed-up matches half-approx coverage in 5x less time
- Demo CLI (`demo/demo.py`) for one-command end-to-end runs

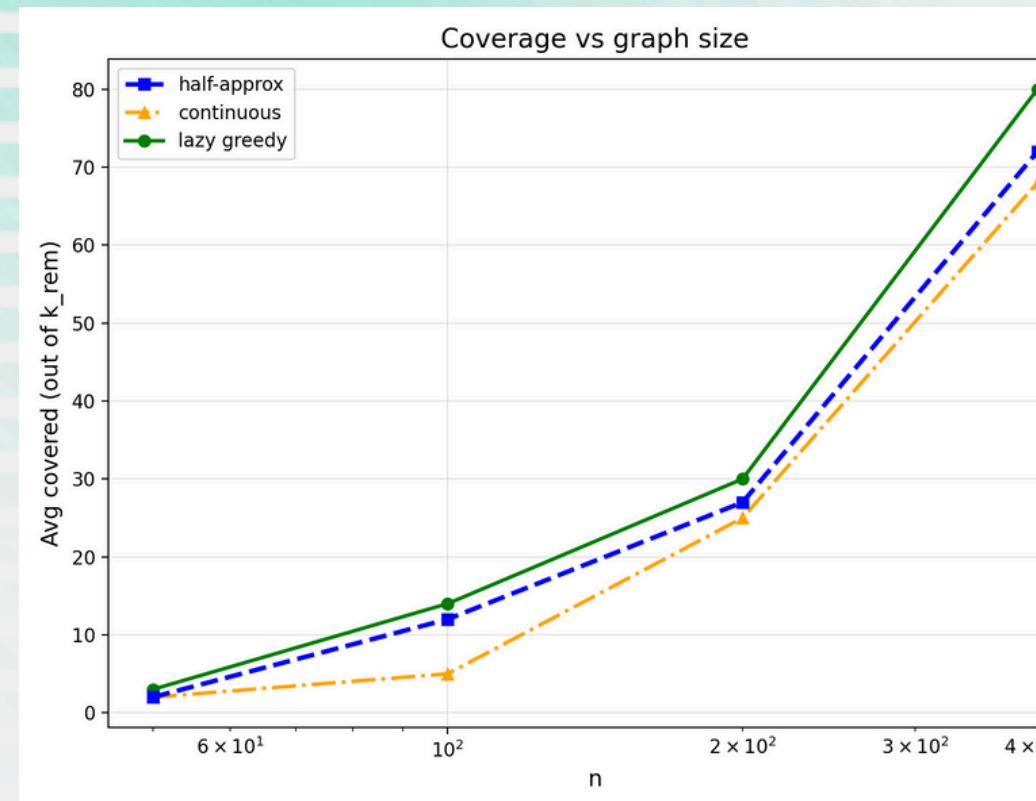
Results

Coverage vs. Runtime (ER graph: $n=500$, $p=0.02$, $k=0.6t$, $D^*=3$)

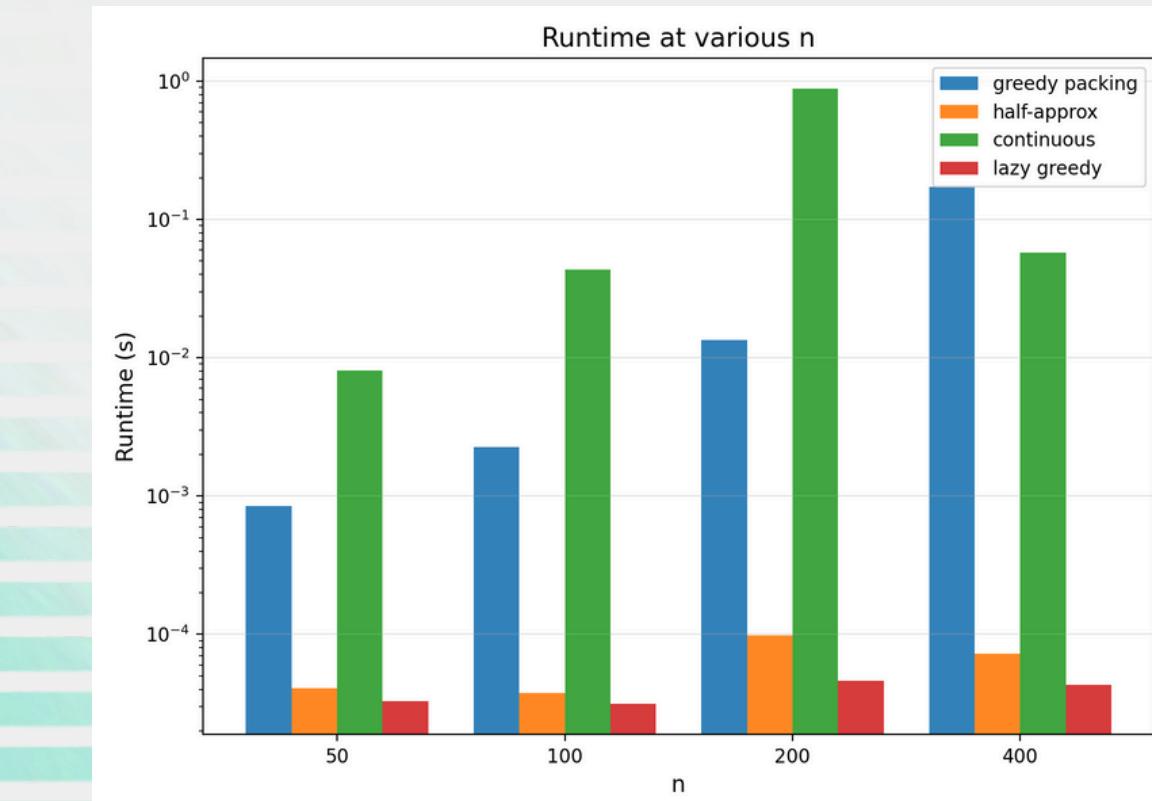
- Half-approx: covers ~45 % of the remaining k in ~0.02 s
- Continuous-greedy: covers ~65 % in ~0.80 s
- Lazy-greedy: matches half-approx coverage (~45 %) in ~0.03 s

Broadcast Rounds

- Full pipeline (greedy + PMCoverlazy) informs k in ~5 rounds
- Greedy-only packing stalls after ~2 rounds

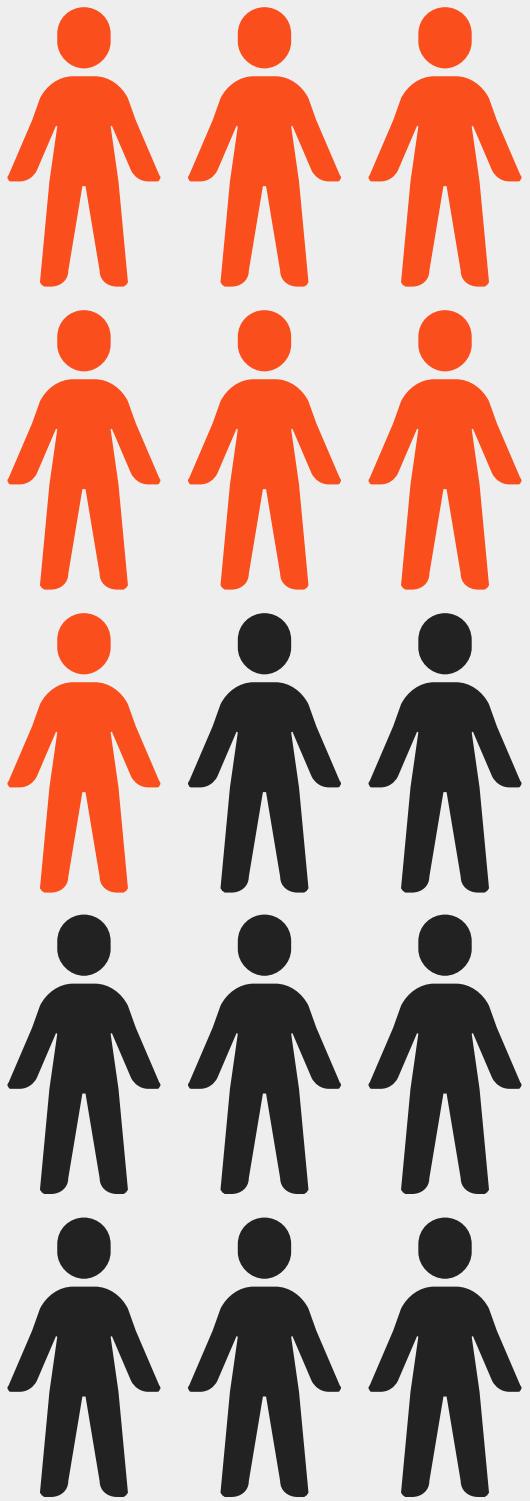


coverage_comparison.png



runtime_comparison.png

Demo



Conclusion

- We delivered a full Python implementation of the directed k-MTM algorithm, matching theoretical guarantees in practice.
- Empirical benchmarks show the trade-off between coverage quality (continuous-greedy) and runtime (lazy-greedy).
- A user-friendly CLI and CI pipeline ensure reproducibility and ease of exploration.

Thank you

PRESENTED BY

Hamza Abdullah