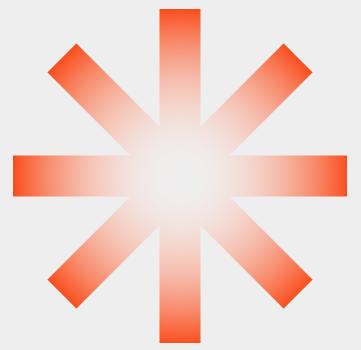


# Directed k-MTM Multicast Implementation

Approximation Algorithms in the Telephone Model

PRESENTED BY

Hamza Abdullah



# Overview

Problem & Motivation

Results

Algorithm

Demo

Implementation

Conclusion

# Problem & Motivation

- Objective: Inform any  $k \leq t$  terminals in a directed network under the telephone model
  - Each informed node may call one neighbor per synchronous round
- Why it matters:
  - Database replication consistency
  - Sensor-network firmware updates
  - Vector-clock synchronization delay
- Core challenge:
  - NP-hard combinatorial problem
  - Matching constraint (one call/node/round) limits throughput

# Algorithm

- Greedy Packing
  - Extract  $\rho = \text{sq\_root}[k]$  vertex-disjoint subtrees of depth  $\leq D^*$ , each covering  $\rho$  terminals.
- Partition Matroid Cover
  - Cover the remaining  $k - \lfloor \rho \rfloor$  terminals via:
    - Half-approx greedy ( $12/\sqrt{2} \approx 12.7$ )
    - Continuous-greedy ( $(1 - 1/e) \cdot \sqrt{e} \approx 1.17$ )
    - Lazy-greedy speed-up
- Stitch & Simulate
  - Combine subtrees + cover edges into one tree, then simulate telephone-model rounds to count broadcast time.

# Implementation

- `graph_loader.py` – generates ER and clique graphs, picks  $t_{\text{ttt}}$  and  $k_{\text{kkk}}$  ratios
- `greedy_packing.py` – extracts the  $\lceil k \rceil \lfloor \log_2 k \rfloor + \lceil \sqrt{k} \rceil$  depth- $D \times D^*D \times$  subtrees
- `pmcover.py`, `pmcover_continuous.py`, `pmcover_lazy.py` – three variants of matroid-constrained set-cover
- `complete.py` – stitches subtrees and cover edges into a single multicast tree
- `simulator.py` – runs the telephone-model broadcast rounds on the final tree

## Enhancements:

- Lazy-greedy speed-up matches half-approx coverage in  $\sim 5 \times$  less time
- Demo CLI (`demo/demo.py`) for one-command end-to-end runs

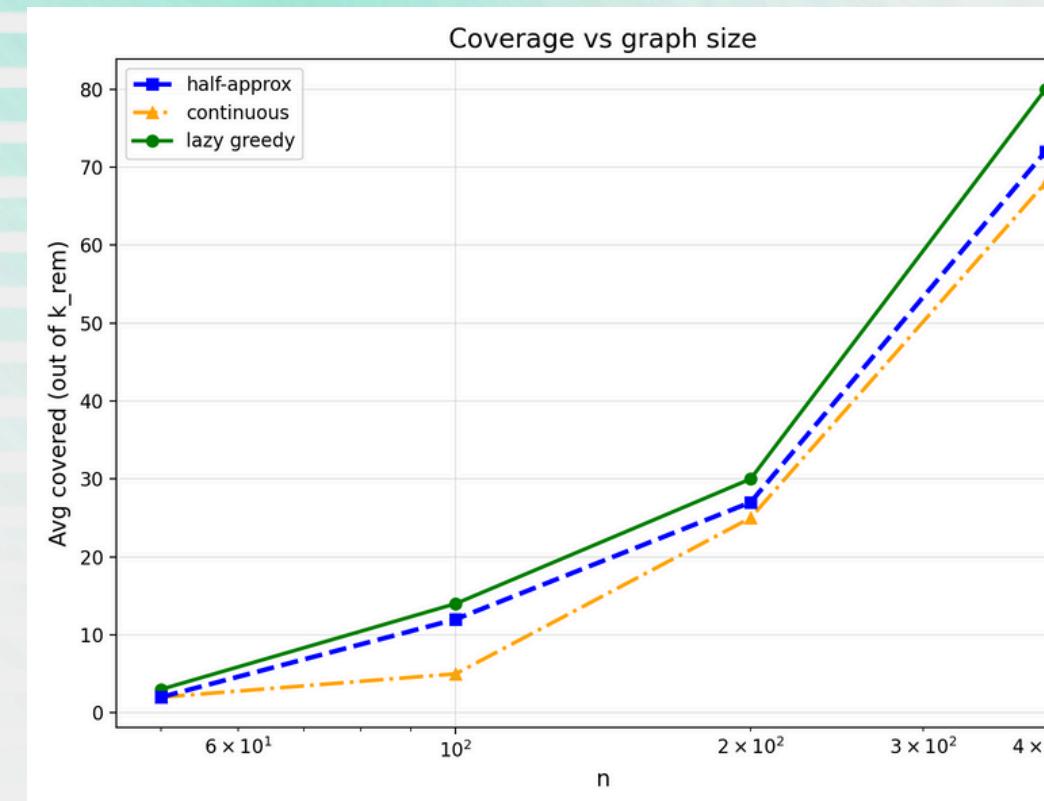
# Results

Coverage vs. Runtime (ER graph:  $n=500$ ,  $p=0.02$ ,  $k=0.6t$ ,  $D^*=3$ )

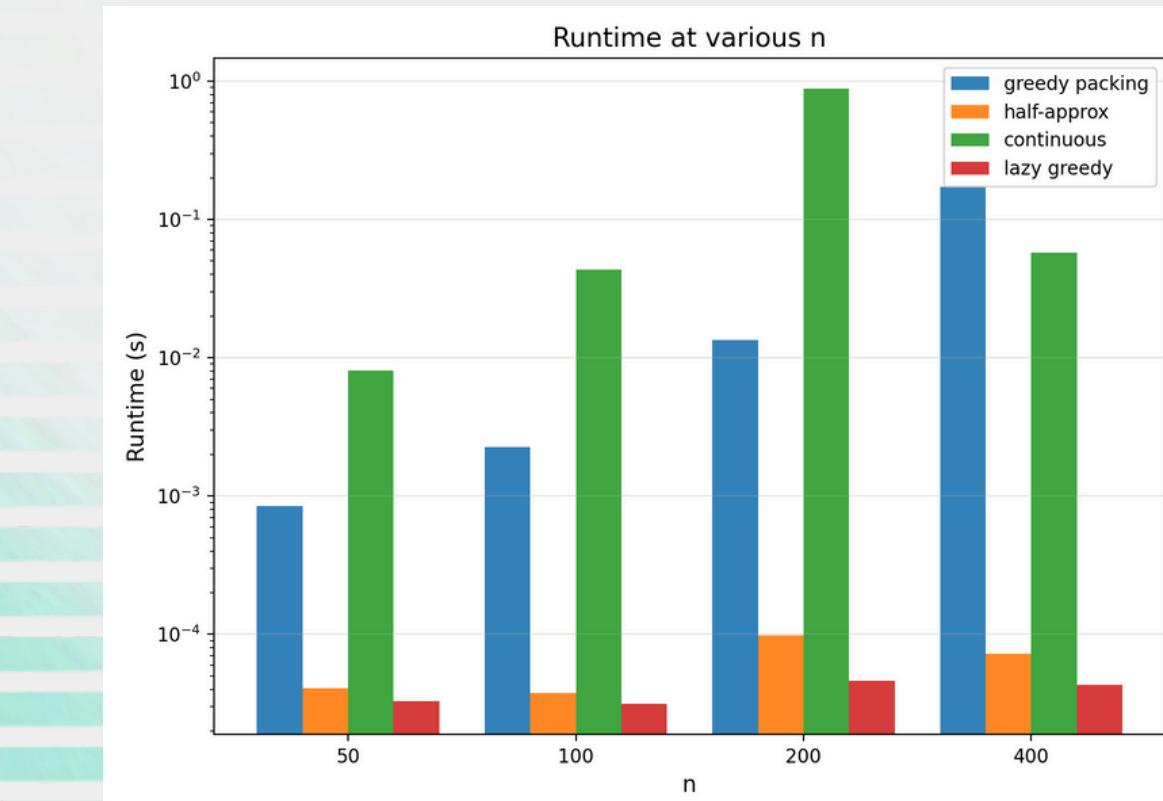
- Half-approx: covers ~45 % of the remaining  $k_{\text{rem}}$  in ~0.02 s
- Continuous-greedy: covers ~65 % in ~0.80 s
- Lazy-greedy: matches half-approx coverage (~45 %) in ~0.03 s

Broadcast Rounds

- Full pipeline (greedy + PMCoverlazy) informs kkk in ~5 rounds
- Greedy-only packing stalls after ~2 rounds (covers only  $\rho_2 \rho_1^2 \rho_2$  terminals)

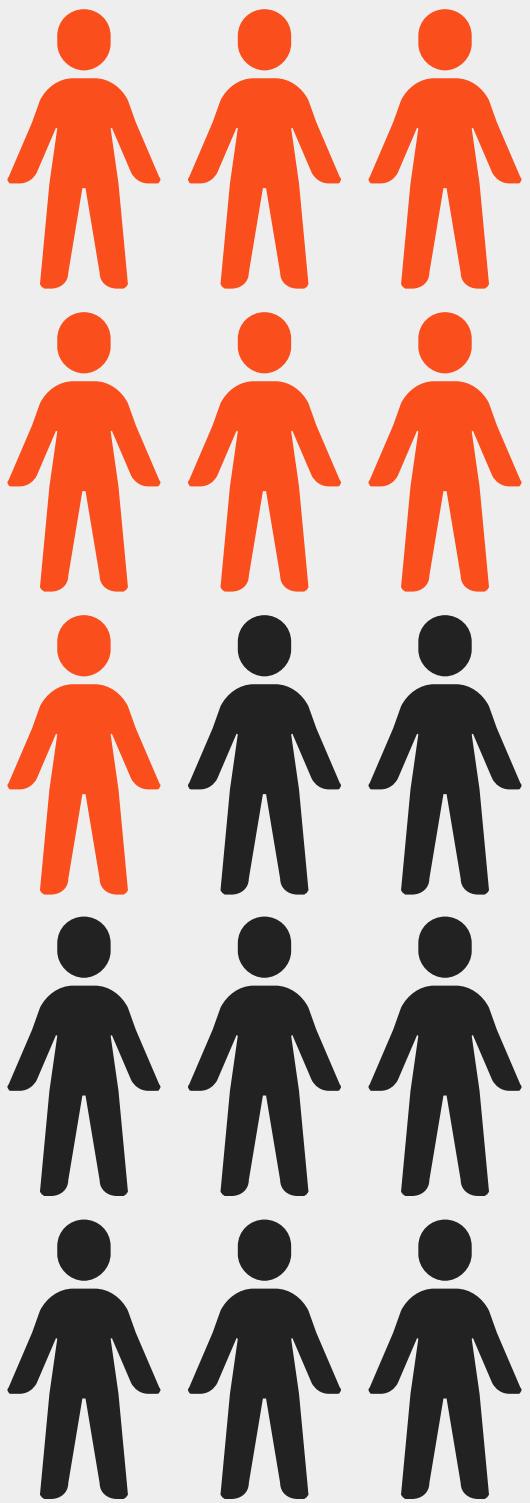


coverage\_comparison.png



runtime\_comparison.png

# Demo



# Conclusion

- We delivered a full Python implementation of the directed k-MTM algorithm, matching theoretical guarantees in practice.
- Empirical benchmarks show the trade-off between coverage quality (continuous-greedy) and runtime (lazy-greedy).
- A user-friendly CLI and CI pipeline ensure reproducibility and ease of exploration.

# Conclusion

Briefly elaborate what you want to discuss.

**Paragraph 1** Present with ease and wow any audience with Canva Presentations. Choose from over a thousand professionally-made templates to fit any objective or topic. Make it your own by customizing it with text and photos.

**Paragraph 2** Present with ease and wow any audience with Canva Presentations. Choose from over a thousand professionally-made templates to fit any objective or topic. Make it your own by customizing it with text and photos.

# Thank you

PRESENTED BY

Hamza Abdullah